

Group B: TJ Gardner-Souza, Amanda Huang, Luc Ravenelle, Erik Wilder, and Muskan Yadav
Professor Ding
MATH 7343
April 25, 2022

A Statistical Analysis of Heart Disease & Stroke

Objective

Heart Disease and Stroke are some of the leading causes of death in America. We investigated an open-source stroke dataset from Kaggle that contains information on various patients' health metrics and whether or not they suffered a stroke. The objective was to find variables that were correlated with a patient who had suffered a stroke, and to determine if an accurate classification could be built to predict whether a new patient would be likely to have a stroke or not.

Study Design

Target Population: Global human population (all ages, with no distinct region specified)

Sampled Population: Those patients with data recorded (unclear, source of data is confidential)

Sample: 5,111 people

Parameter in question: Proportion of patients who have experienced stroke

The dataset chosen for analysis involves patients from an unknown source, regardless the dataset is intended for a global population. People of all ages with various pre-existing health conditions and statuses have been recorded in this dataset with a final, exhaustive parameter indicating their stroke diagnosis. According to the World Health Organization (WHO) stroke is the second largest cause of death among people worldwide, with its share standing at about eleven percent of deaths each year. This dataset's purpose and our goal is to find other health status input parameters that may be strong predictors of stroke, and additionally, make a prediction for a patient.

The dataset hosts a sample of 5,111 patients of unknown origin or location. Although the sample is large, we must call attention to the fact that randomness cannot be guaranteed, because the methods of collection and source are unknown. It must also be stated that this study is observational in nature, as the parameters and their respective statistics cannot be collected in a proper or controlled means. Therefore, the outcome of this analysis will only show association and not make any causal conclusions. All variables that may or may not show statistical significance in our analysis will infer an association or non-association between themselves and stroke.

Data Analysis

The dataset was imported from the Kaggle repository. It contains 10 predictor variables (7 categorical and 3 numerical) and a binary variable indicating whether or not the patient had a stroke. To prepare the data, null values in the predictor variables were replaced so that R could understand the data set. We then normalized the numerical data such as glucose, BMI, and age. The categorical variables include gender, hypertension, heart disease, work type, urban versus rural residence, smoking status, and stroke history. For regression formatting, we created ‘dummy’ variables from the categorical variables and up-sampled the minority class.

We used contingency tables to determine the frequency of an individual categorical variable against stroke; specifically, stroke versus hypertension, stroke versus heart disease, and stroke versus smoking status. Results of the contingency tests can be pictured below:

```
      stroke
hypertension 0    1
0 4309 149
1  391  60
      stroke
hypertension 0    1
0 0.96657694 0.03342306
1 0.86696231 0.13303769
```

Pearson's Chi-squared test

data: cont_table_hyp

X-squared = 99.704, df = 1, p-value < 2.2e-16

```
      stroke
heart_disease 0    1
0 4497 169
1  203  40
      stroke
heart_disease 0    1
0 0.96378054 0.03621946
1 0.83539095 0.16460905
```

Pearson's Chi-squared test

data: cont_table_hd

X-squared = 93.403, df = 1, p-value < 2.2e-16

```
      stroke
smoking 0    1
formerly smoked 780 57
never smoked 1768 84
smokes 698 39
Unknown 1454 29
      stroke
smoking 0    1
formerly smoked 0.93189964 0.06810036
never smoked 0.95464363 0.04535637
smokes 0.94708277 0.05291723
Unknown 0.98044504 0.01955496
```

Pearson's Chi-squared test

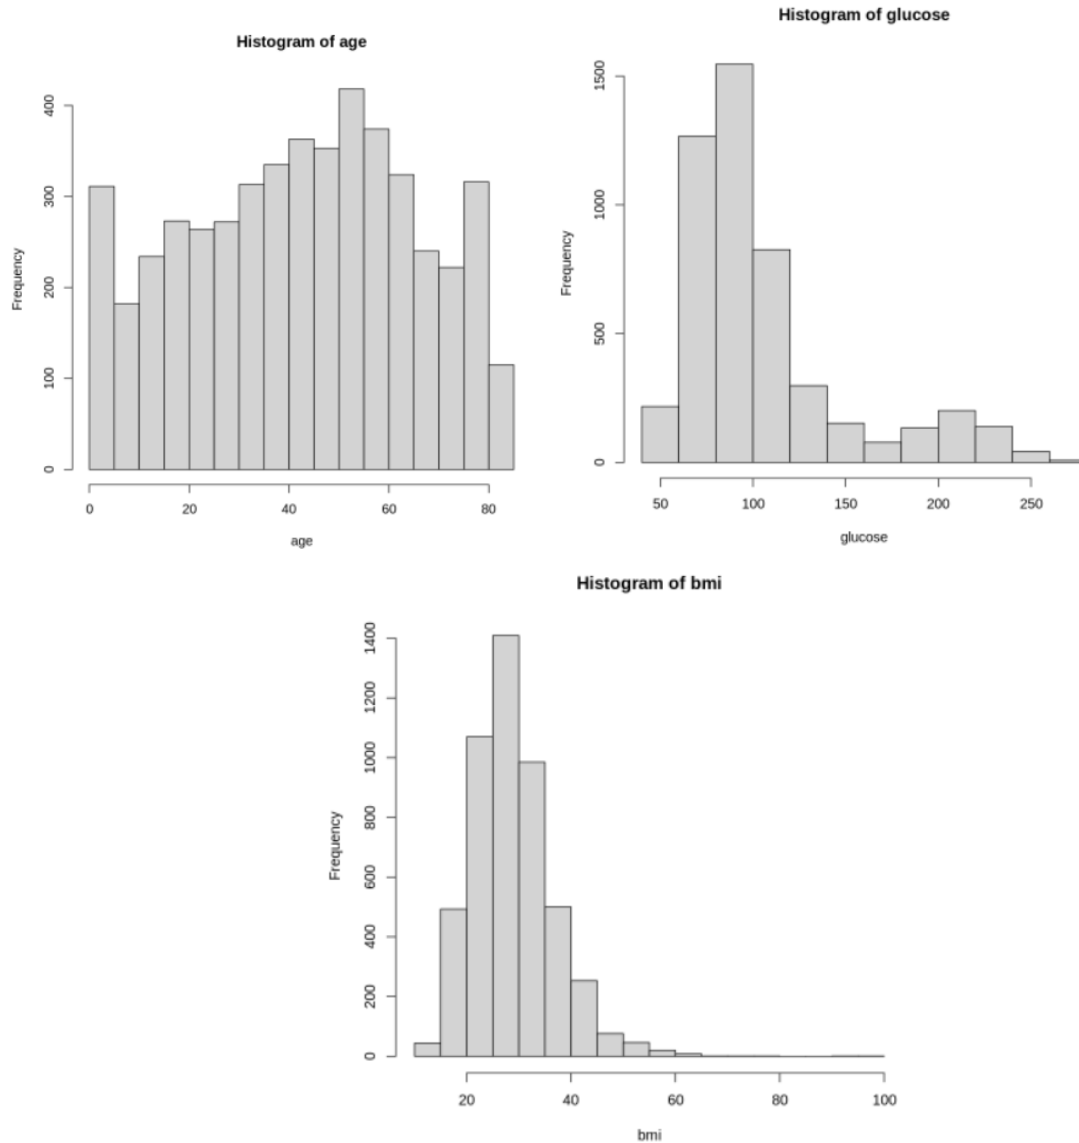
data: cont_table_smoke

X-squared = 34.943, df = 3, p-value = 1.252e-07

For each of the Chi-square tests performed above, the null hypothesis is that there is no relationship between the predictor variables (hypertension, heart disease, and smoking status) and the target variable (stroke). If this were the case, we would see the same distribution of stroke vs. non-stroke patients across the different categories of the independent variables. However, in this case, we see that the p-value for each of the tests performed is less than the level of significance, $\alpha = 0.05$. Therefore, we can reject the null hypothesis that no relationship

exists between each of the predictor variables and the target variable. Thus, we conclude that there is sufficient evidence to suggest that hypertension, heart disease, and smoking status are all correlated with stroke.

Before performing tests on the numerical variables, histograms were implemented to categorize the type of distribution and to determine the correct type of test that should be executed. The outputs of the histograms received for age, glucose and BMI are pictured below:



Normally distributed data looks similar to a bell shaped curve with one peak and is symmetric around the mean. Based on the histograms above, age, glucose and BMI are not normally distributed. Thus, warranting the Wilcoxon rank-sum test.

The Wilcoxon rank-sum is a nonparametric test for unpaired data that can be used to compare two samples from independent populations and does not require for the underlying populations to be normally distributed or that their variances be equal. For each of the Wilcoxon tests – comparing stroke to numeric data: age, average glucose level, and BMI– the results return p-values that are significantly less than the level of significance, $\alpha = 0.05$. With the p-values less than the level of significance for all three tests, we concluded that we reject our null hypothesis for these tests and conclude that there is evidence that the medians of the two tested populations are not equal. The outputs of the Wilcoxon tests are pictured below:

```
Wilcoxon rank sum test with continuity correction

data:  stroke$age and nonstroke$age
W = 821280, p-value < 2.2e-16
alternative hypothesis: true location shift is not equal to 0

Wilcoxon rank sum test with continuity correction

data:  stroke$avg_glucose_level and nonstroke$avg_glucose_level
W = 614278, p-value = 8.168e-10
alternative hypothesis: true location shift is not equal to 0

Wilcoxon rank sum test with continuity correction

data:  stroke$bmi and nonstroke$bmi
W = 569022, p-value = 0.0001026
alternative hypothesis: true location shift is not equal to 0
```

Results

To determine if it is possible to build an accurate classifier for identifying a new patient as “at-risk” for having a stroke, we have performed a series of logistic regressions. Our first attempt used all the variables and split the data into 75% training and 25% testing data and used all variables. Numerical data were normalized and categorical data were placed into ‘dummy’ columns to achieve the best results. The ‘Formerly Smoked’ Category had only the ‘0’ class, so this variable was omitted from all regression attempts. The prediction result of this model is pictured below:

```
stroke.pred    0    1
              0 1175   53
```

Due to the imbalance of stroke patients and non-stroke patients, the model did not make any ‘stroke’ predictions. To troubleshoot this, we used an over-sampling technique to increase the size of the minority class. Another logistic regression was performed with the same input variables and a balanced target class. The results are as given below:

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.7603	-0.7749	-0.2240	0.8107	2.6343

Coefficients: (4 not defined because of singularities)

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-12.01092	198.36696	-0.061	0.951718
male1	-0.01136	0.08587	-0.132	0.894740
female1	NA	NA	NA	NA
age	1.36938	0.06518	21.010	< 2e-16 ***
hypertension1	0.77725	0.11602	6.699	2.10e-11 ***
heart_disease1	0.74617	0.14513	5.141	2.73e-07 ***
married1	0.21629	0.13288	1.628	0.103591
unmarried1	NA	NA	NA	NA
private1	10.58041	198.36696	0.053	0.957463
govt1	10.27298	198.36699	0.052	0.958698
self1	10.61543	198.36698	0.054	0.957322
children1	11.36336	198.36714	0.057	0.954319
never_worked1	NA	NA	NA	NA
urban1	0.13695	0.08214	1.667	0.095467 .
rural1	NA	NA	NA	NA
glucose	0.13422	0.03241	4.141	3.46e-05 ***
bmi	0.16358	0.04639	3.526	0.000421 ***
smokes1	0.34374	0.12596	2.729	0.006354 **
never_smoked1	-0.13400	0.10622	-1.262	0.207118
unknown_smoker1	-0.07198	0.12930	-0.557	0.577760

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 5100.1 on 3680 degrees of freedom
Residual deviance: 3661.7 on 3665 degrees of freedom
AIC: 3693.7

Number of Fisher Scoring iterations: 12

This logistic regression identifies six variables to be correlated with stroke. Age, hypertension, heart disease, glucose, BMI, and smokes all had p-values less than the level of significance, $\alpha = 0.05$. These results are in line with our previous Chi-square and Wilcoxon tests, so we are confident there exists a correlation between these six variables and stroke. Classification results are as shown below:

```
stroke.pred  0  1
            0 989 18
            1 186 35
```

```
Call:
accuracy.meas(response = test$stroke, predicted = balanced_result)
```

Examples are labelled as positive when predicted is greater than 0.5

```
precision: 0.158
recall: 0.660
F: 0.128
```

This model has a precision of 0.158, which can be interpreted as that around 85% of positive classifications were incorrect. The recall is 0.660, meaning that the model classified 66% of true positives correctly. The F-score is an overall accuracy measure for the model, which in our case is low, suggesting that this model does not perform well overall.

As a measure to improve the performance, we omitted variables that were not found to be significant. Results from this model are as pictured below:

```
Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -1.27066    0.06496  -19.561  < 2e-16 ***
age           1.36228    0.05748   23.701  < 2e-16 ***
hypertension1  0.77463    0.11273    6.872 6.35e-12 ***
heart_disease1 0.74139    0.14348    5.167 2.38e-07 ***
glucose       0.13898    0.03195    4.349 1.37e-05 ***
bmi           0.14997    0.04503    3.331 0.000866 ***
smokes1       0.41448    0.10549    3.929 8.53e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 5100.1 on 3680 degrees of freedom
Residual deviance: 3681.6 on 3674 degrees of freedom
AIC: 3695.6
```

Number of Fisher Scoring iterations: 5


```
stroke.pred  0  1
            0 994 16
            1 181 37
```

```
Call:
accuracy.meas(response = test$stroke, predicted = balanced_result_refined)
```

Examples are labelled as positive when predicted is greater than 0.5

```
precision: 0.170
recall: 0.698
F: 0.137
```

All variables have a p-value less than the level of significance, $\alpha = 0.05$, and hence all of them are correlated with stroke, as was expected. We observe slight improvements in precision, accuracy, and recall as well.



It is noticed that none of the numerical variables show a high correlation with each other, so we do not need to remove any of these from the model. Additionally, to test correlation among the categorical variables, we performed a series of chi-square tests among them:

heart_disease		
hypertension	0	1
0	4273	185
1	393	58
heart_disease		
hypertension	0	1
0	0.95850157	0.04149843
1	0.87139690	0.12860310

Pearson's Chi-squared test

data: cont_table_hyp

X-squared = 66.045, df = 1, p-value = 4.407e-16

smoking						
heart_disease	formerly	smoked	never	smoked	smokes	Unknown
0		767		1771	682	1446
1		70		81	55	37
smoking						
heart_disease	formerly	smoked	never	smoked	smokes	Unknown
0	0.1643806	0.3795542	0.1461637	0.3099014		
1	0.2880658	0.3333333	0.2263374	0.1522634		

Pearson's Chi-squared test

data: cont_table_hd

X-squared = 50.919, df = 3, p-value = 5.09e-11

hypertension		
smoking	0	1
formerly smoked	727	110
never smoked	1636	216
smokes	655	82
Unknown	1440	43
hypertension		
smoking	0	1
formerly smoked	0.86857826	0.13142174
never smoked	0.88336933	0.11663067
smokes	0.88873813	0.11126187
Unknown	0.97100472	0.02899528

Pearson's Chi-squared test

data: cont_table_smoke

X-squared = 102.89, df = 3, p-value < 2.2e-16

We consider that the null hypothesis for each of these tests is that the variables are not correlated. Since each test has a p-value less than the level of significance, $\alpha = 0.05$, we reject the assumed null hypothesis that these variables are not correlated. There is sufficient evidence to suggest that a correlation exists amongst all of the categorical predictor variables. For our final regression model, we kept all numerical variables but only selected the most significant categorical predictor value, hypertension. The results from the final model are as presented below:

```
Coefficients:
      Estimate Std. Error z value Pr(>|z|)
(Intercept)  -1.14407    0.05960 -19.196 < 2e-16 ***
age           1.39567    0.05596  24.940 < 2e-16 ***
hypertension1 0.78408    0.11225   6.985 2.84e-12 ***
glucose       0.16784    0.03137   5.350 8.80e-08 ***
bmi           0.15564    0.04483   3.472 0.000517 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 5100.1  on 3680  degrees of freedom
Residual deviance: 3736.7  on 3676  degrees of freedom
AIC: 3746.7

Number of Fisher Scoring iterations: 5

stroke.pred   0    1
             0 987  16
             1 188  37

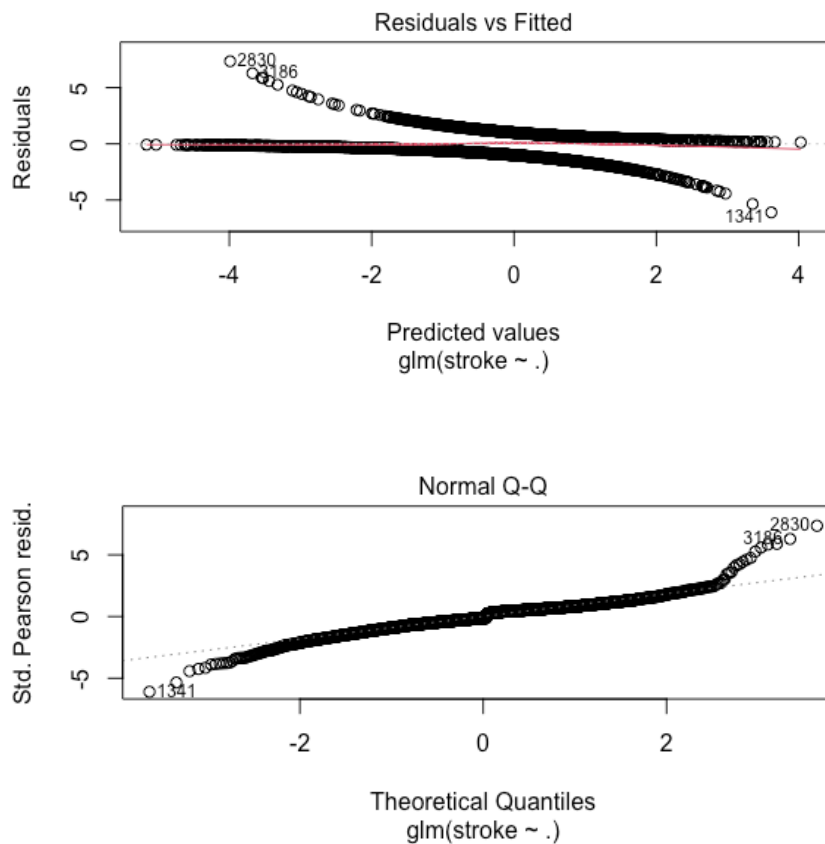
Call:
accuracy.meas(response = test$stroke, predicted = balanced_result_correlated)

Examples are labelled as positive when predicted is greater than 0.5

precision: 0.164
recall: 0.698
F: 0.133
```

Removing correlated categorical features did not result in a significant improvement. However since this model is seen to be the least complex, we select this to be the best model. It was noted that 16.4% of positive classifications were correct, and 70% of true positives were identified correctly. The F-score of 0.133 suggests that this model does not generally perform well.

To add to the analysis of the model's performance, we can plot the residuals. R will give several plot options in which we can examine the errors of our model. Here, we have included a plot of the residuals plotted against the fitted values, as well as a plot that shows Pearson residuals plotted with the theoretical quantiles. Examining both plots, we can further our conclusion about the poor performance of the model mentioned previously. In the residual vs fitted values plot, we can see a pretty large gap between the two values, as well as some pretty unrealistically large values. Even though it may be possible mathematically, it does not make sense to have values outside of 0 and 1, as it is a logistic model.



Examining residuals is an effective way to identify incorrect inferences in data and possible peculiarities that could lead us there. One peculiarity shown in our residual plots is that both show several data points that are classified as outliers. These points are shown to be outside of the expected range of possible values and could be impacting our model's efficacy. In this case, an experienced statistician would simply remove them in order to improve the prediction of the model. However, the removal needs to be justified. Upon examining our outliers, we concluded that removal of these outlier data points would be inappropriate and compromise the model. Let's take the data point #1341, for example. This patient has a fairly high average glucose level of 178.76 with a below average BMI of 24.1. This patient registers as an outlier in our residual plot, but the data makes sense. The patient is a middle aged male, with the preexisting condition of heart disease. After examining a few of these outliers for possible removal, we decided that it did not make sense to do so.

Discussion

The goal of this project was to identify variables that were correlated with stroke and investigate if it is possible to build an accurate classifier to categorize a patient as at-risk for having a stroke. The Chi-square tests and Wilcoxon tests showed that age, heart disease, hypertension, smoking status, BMI, and average glucose level are all correlated. However, we weren't able to build an accurate classifier to prove the same.